

## SUMMARY

---

- Over 8 years of experience including 5+ years of experience with Hadoop Ecosystem in installation and administrated of all UNIX/LINUX servers and configuration of different Hadoop eco-system components in the existing cluster project.
- Experience in set up and installation of Hadoop Cluster and additional components.
- Design and set up security for Hadoop cluster using Kerberos, Ranger, Splunk etc.
- Experience in set up High Available Cluster and Disaster Recovery cluster.
- Experience in Cluster planning and capacity planning of the cluster.
- Experience in set up Hive, Oozie, Spark, Sqoop, Kafka, Custom Components etc.
- Set up Name Node High Availability, RM, Hive services.
- Experience in set up Data Lake Architecture like Landing Zone, Data Lake and Reservoir in HDFS layer.
- Designed process for Onboarding users and Projects into Hadoop cluster.
- Experience in enabling performance tuning for MR, Hive, Spark, HBase, Kafka components.
- Set up the secure connectivity between source and Hadoop as well Hadoop to third party tools like tableau.
- Enabled the HDFS encryption in Hadoop cluster.
- Set up ranger policies for the secure access.
- Enable the end to end monitoring for Hadoop cluster includes Linux, OS and Hadoop monitoring using Nagios and Ganglia tools.
- Set up L1, L2 and L3 support using ITIL process and publish the escalation matrix to all the users.
- Developed data ingestion scripts using Sqoop, Nifi, Kafka.
- Set up JDBC and ODBC connectivity for the external applications.
- Developed data processing jobs using pig, spark and Map Reduce technologies.
- Experience on Oozie workflow development and deployment.
- Experience on DevOps deployment model using VSTS.
- Involved in project planning and Agile methodologies.
- Set up process for Code review prior to deploy the code in PROD environment.

## TECHNICAL SKILLS

---

Hadoop Distribution:	Cloudera Distribution of Hadoop (CDH), HDP, HDF
Scripting:	Java, PL SQL, Shell Script, Perl, Python
Big Data Tools:	MapReduce, HDFS, YARN, Hive, Pig, Sqoop, Oozie, Kafka, Hortonworks, Ambari, Knox, Phoenix, Impala, Storm, Nifi.
Tools:	SQL Server Management, Studio, Jenkins, GitHub, Ranger Test NG, Junit, Hcatalog, Tableau, GMS, BMC Remedy, Eclipse, Toad,
Databases:	MySQL, NoSQL, Couchbase, Influx DB, Teradata, HBase, Hive, Cassandra, Oracle.
Processes:	ITIL, Systems Administration, Incident Management, Release Management Change Management.

## PROFESSIONAL WORK EXPERIENCE:

---

**Designation:** Engineer – Hadoop Administrator.

**Duration:** From September 01, 2017 to Till now.

**Client:** AT&T Alpharetta.

### Responsibilities: -

- Design/Install/Configure Cloudera cluster with HDFS, Yarn, MapReduce, Spark, Zookeeper, Impala, HBase services
- Design and develop platform architecture and provide solutions to the projects.
- Discussed with Project teams and validate project architecture and provide suitable solutions.
- Set up capacity queue for all projects based on the utilization, dynamic queue creation.
- Resolve customer & project architect clarifications and change architecture based the cluster set up.
- Change the configurations based on the customer needs, without effecting to the other projects.
- Implemented Capacity Queue Management, Quota Management successfully.
- Provided L2 and L3 support for critical issues.
- Install Network SSL certifications both source and Hadoop side.
- Implemented SMARTSENSE recommendations for cluster performance tuning.
- Continuous monitoring and managing the Hadoop cluster using Cloudera Manager.
- Worked on fixing the cluster issues and Configuring High Availability for Name Node
- Developed shell script for edge node clean up and logs clean up.
- Design best practices for each service in the cluster and published to projects teams.
- Professionalize the projects and successfully deploy the applications in Hadoop cluster.
- Designed and developed data ingest factory model to ingest data very easily with minimal set ups.

### Environment:

Hadoop, Cloudera Manager, MapReduce, Yarn, Impala, Pig, Sqoop, Oozie, Flume, Zookeeper, HBase.

**Designation:** Senior Operational Specialist.

**Duration:** From April 01, 2016 to August 24, 2017.

**Client:** IBM Hyderabad

### Responsibilities: -

- Experience in configuring, installing and managing MapR, Hortonworks & Cloudera, Hive and Spark.
- Skilled in scheduling recurring Pig and Hive jobs using Rundeck.
- Deployed Grafana Dashboards for monitoring cluster nodes using Graphite as a Data Source and collect as a metric sender. Used Pig as ETL tool to do transformations, event joins, filter and some pre-aggregations.
- Hadoop Ecosystem Cloudera, Hortonworks, Hadoop, MapR, HDFS, HBase, Yarn, Zookeeper, Nagios, Hive, Pig, Ambari Spark Impala. Installed and configured Drill, Fuse and Impala on MapR-5.1.
- Maintaining the Operations, installations, configuration of 150+ node clusters with MapR distribution.
- Implementation of Kerberized Hadoop Ecosystem. Using Sqoop and Nifi in a Kerberized system to transfer data from relational databases like MySQL to HDFS.
- Managed mission-critical Hadoop cluster and Kafka at production scale, especially Cloudera distribution.
- Hands on working experience with Devops tools, chef, puppet, Jenkins, git, maven, Ansible.
- Installation, Upgrade, Configuration of Monitoring Tools (MySQL Enterprise Monitor, New Relic and DataDog APM monitoring). Experienced in MapR cluster to monitoring through ITRS.
- Implemented Cloudera Impala on top of hive for faster querying for user.
- Wrote workflows which include data cleansing Pig actions and hive actions.
- Developed Spark SQL to load tables into HDFS to run select queries on top.

- Created Hive tables on top of HDFS files and designed queries to run on top.
- Extended Hive and Pig core functionality by designing custom UDFs.
- Experience on DNS, NFS, and DHCP, printing, mail, web, and FTP services for the enterprise.
- Experience on Manages UNIX account maintenance including additions, changes, and removals.

**Environment:**

Kerberos, Linux Admin, Kafka, YARN, Spark, HBase, Hive, Impala, SOLR, Java Hadoop cluster, HDFS, Ambari, Ganglia, Nagios, Cloudera, MapR.

**Designation:** Staff Associate-Technical Support.

**Duration:** From May 14, 2012 to March 31, 2016.

**Client:** AT&T Hyderabad

**Responsibilities: -**

- Assist in Install and configuration of Hive, Pig, Sqoop, Flume, Oozie and HBase on the Hadoop cluster with latest patches.
- Implemented Flow file, Connections, Flow controller and Process Group as part of Nifi Process for automating the movement of data.
- Used Spark Streaming to divide streaming data into batches as an input to spark engine for batch processing. Mentored EQM team for creating Hive queries to test use cases.
- Installed/Configured/Maintained Apache Hadoop clusters for application development and Hadoop tools like Hive, Pig, HBase, Zookeeper and Sqoop.
- Involved in installation and configuration of Kerberos security setup on CDH5.5 cluster.
- Involved in installation and configuration of LDAP server and integrated with Kerberos on cluster.
- Worked with Sentry configuration to provide centralized security to Hadoop services.
- Monitor critical services and provide on call support to the production team on various issues.
- Assist in Install and configuration of Hive, Pig, Sqoop, Flume, Oozie and HBase on the Hadoop cluster with latest patches.
- Involved in performance tuning of various Hadoop ecosystem components like YARN, MRv2.
- Implemented the Kerberos security software to CDH cluster for user level as well as service level to provide strong security to the cluster.
- Troubleshooting, diagnosing, tuning, and solving Hadoop issues.
- Continuous monitoring and managing the Hadoop cluster using Cloudera Manager.
- Commissioning and decommissioning the nodes across cluster.
- Creating volumes and snapshots through the MAPR CONTROL SYSTEM. Experience in support of our MapR Hadoop installations in Dev, QA and Product also.
- Install various Splunk Applications such as Cisco for Splunk, Windows for Splunk, and VMware for Splunk. Created MapR DB tables and involved in loading data into those tables.
- Expertise in Virtualizations System Administration of VMware, VMware Server, VMware Lab Manager, Cloud, Amazon EC2 & S3 web services.
- Responsible for developing data pipeline using HD Insight, flume, Sqoop and pig to extract the data from weblogs and store in HDFS.
- Experience in understanding Hadoop multiple data processing engines such as interactive SQL, real time streaming, data science and batch processing to handle data stored in a single platform in Yarn.
- Working with data delivery teams to setup new Hadoop users.
- This job includes setting up Linux users, setting up Kerberos principals and testing HDFS, Hive.
- Install OS and administrated Hadoop stack with CDH5 (with YARN) Cloudera Distribution including configuration management, monitoring, debugging, and performance tuning Scripting Hadoop package installation and configuration to support fully automated deployments.

- Worked with systems engineering team to plan and deploy new Hadoop environments and expand existing Hadoop clusters.
- Installed and configured opensource Puppet server for firmwide configuration management.
- Work at optimizing volumes and EC2 instances and created multiple VPC instances.
- Create custom sized VPC, subnets, NAT to ensure successful deployment of Web applications and database templates.
- Investigated customer LDAP and LDAPS configurations to ensure certificates and Active Directory were correctly configured.
- Experience installing and configuring SSH (Secure Shell) encryption in order to access securely on Ubuntu and Red hat Linux.
- Worked on troubleshooting for LDAP and SiteMinder issues with Support Teams for newer initiatives at organization level.
- Management of configuration files, services, packages using puppet.
- Maintain Jenkins continuous integration infrastructure and automated releases to DEV/TEST/STG/PROD environments.
- Implemented a Continuous Delivery framework using Jenkins, Puppet in Linux environment.
- Involved in complete Administration tasks on UNIX, Red Hat Linux and Solaris and documentation for the projects executed.
- Setup multipath and LVM volumes on Linux servers.
- Writing custom Chef Recipes for managing configurations and dynamic automation processes for production systems and zenoss dev environments.
- Troubleshoot various OS and network related issues and provided RCA.
- Worked on sev1,2,3 tickets, service requests and production changes and releases.
- Install, configure and maintain web server like Apache Web Server.
- Streamlined installation of OpenShift on partner cloud infrastructure such as AWS and Azure.
- Create policies for projects according to client requirements.
- Experience in backup/restore of PostgreSQL databases.
- Created host clusters for high availability of VMs with vMotion capabilities.

**Environment:-** Hive, Pig, HBase, Zookeeper and Sqoop, ETL, Ambari 2.0, HDP , Linux Cent OS, MongoDB, Cassandra, Ganglia and Cloudera Manager , AWS, S3, EBS, REHL 6/7, Elastic Load balancer (ELB), PostgreSQL, Auto Scaling groups, VPC, IAM, Puppet, UNIX, MongoDB, ETL ,Cloud Watch, Informix 11.5, Glacier, Solaris 10, bash scripts, PHP script and Linux environment.

**Designation:** Staff Associate-Technical Support.

**Duration:** From November 15, 2011 to May 14, 2012

**Client:** AT&T Hyderabad

- Provided 24x7 on-call supports in debugging and fixing issues related to Linux, Solaris, HP-UX Installation/Maintenance of Hardware/Software in Production, Development & Test Environment as an integral part of the Unix/Linux (RHEL/SUSE/SOLARIS/HP-UX/AIX) Support team.
- Installation Red hat Linux Enterprise Server 5/6 on Dell and HP x86 HW.
- Planning and implementing Backup and Restore procedures using Ufsdump, Ufsrestore, Tar" and "Cpio".
- Installed and configured the Red Hat Linux 5.1 on HP-DI585 servers using Kick Start.
- Monitoring day-to-day administration and maintenance operations of the company network and systems working on Linux and Solaris Systems.
- Responsible for deployment, patching and upgrade of Linux servers in a large datacenter environment.
- Installations of HP Open view, monitoring tool, in servers and worked with monitoring tools such as Nagios and HP Open view.

- Installed and configured the RPM packages using the YUM Software manager.
- Involved in developing custom scripts using Shell (bash, ksh) to automate jobs.
- Defining and Develop plan for Change, Problem & Incident management Process based on ITIL.
- Networking communication skills and protocols such as TCP/IP, Telnet, FTP, NDM, SSH, rlogin.
- Deploying Veritas Clusters and Oracle test databases to implement disaster recovery strategies, ensuring uninterrupted availability of the global systems.
- Also coordinating with storage team and networking teams.
- Responsible for providing 24x7 production support for Linux.
- Automated Kickstart images installation, patching and configuration of 500+ Enterprise Linux servers.
- Built kickstart server for automated Linux server builds.
- Installed Ubuntu servers for migration.
- Created Shell, Bash scripts to automate a variety of tasks.
- NFS and SAN filesystem management - Veritas VxVM, LVM
- Maintained user accounts. Sudo was used for management accounts and faceless. Otherwise the accounts were LDAP.
- Datacenter operations, migration of Linux servers
- Configured the NIS, NIS+ and DNS on Red Hat Linux 5.1 and update NIS maps and organize the RHN Satellite Servers in combination with RHN Proxy Server.
- Open Ldap server & clients, PAM authentication setup on RedHat Linux 6.5/7.1.
- Installed, configured, troubleshoot and maintain Linux Servers and Apache Web server, configuration and maintenance of security and scheduling backups, submitting various types of cron jobs.

**Environment:-**

Red Hat Enterprise Linux 4.x/5.x, Logical Volume Manager for Linux and VMware ESX Server 2.x, Hyper-V Manager  
VMware Vsphere Client, RHEL, Citrix Xen.

# Rama K

- Alpharetta, GA, USA

## Contact Information

- ramhadoop1988@gmail.com (Preferred)
- 4704650474 (Preferred)

## Work History

### Total Work Experience: 2 years

- **Hadoop Admin | at&t**  
Sep 01, 2017 - No End Date

## Skills

- **cluster** | 6yrs | 2019
- **unix** | 6yrs | 2019
- **configuration** | 6yrs | 2019
- **installation** | 6yrs | 2019
- **monitoring** | 6yrs | 2019
- **production** | 6yrs | 2019
- **batch processing** | 5yrs | 2019
- **big data** | 5yrs | 2019
- **capacity management** | 5yrs | 2019
- **cloudera** | 5yrs | 2019
- **high availability** | 5yrs | 2019
- **kerberos** | 5yrs | 2019
- **linux** | 5yrs | 2019
- **performance tuning** | 5yrs | 2019
- **sql** | 5yrs | 2019
- **apache hadoop** | 5yrs | 2019
- **apache hadoop admin** | 5yrs | 2019
- **apache hbase** | 5yrs | 2019
- **apache hive** | 5yrs | 2019

- **apache kafka** | 5yrs | 2019
- **apache solr** | 5yrs | 2019
- **apache spark** | 5yrs | 2019
- **apache sqoop** | 5yrs | 2019
- **apache zookeeper** | 5yrs | 2019
- **chef** | 5yrs | 2019
- **distributed computing** | 5yrs | 2019
- **distribution** | 5yrs | 2019
- **ganglia** | 5yrs | 2019
- **hdfs** | 5yrs | 2019
- **hortonworks** | 5yrs | 2019
- **amazon web services** | 2yrs | 2017
- **backup administration** | 3yrs | 2014
- **apache**
- **hadoop**
- **administration**
- **hbase**
- **hive**
- **solr**

## Work Preferences

- Likely to Switch: Most Likely
- Willing to Relocate: No
- Work Authorization:
  - US
- Work Documents:
  - Have H1
- Security Clearance: No
- Third Party: Yes
- Employment Type:
  - Contract - Corp-to-Corp

## Profile Sources

- Dice:
  - <https://www.dice.com/employer/talent/profile/f6a7de90cc59fe599810568ebf3ba84eb0f1685f>